

# Lecture 4: Model-Free Prediction and Control


Hado van Hasselt

February 1, 2018, UCL

# Background

Sutton & Barto 2018, Chapters 5 + 6

# Recap



- ▶ Reinforcement learning is the science of learning to make decisions
- ▶ Agents can learn a **policy**, **value function** and/or a **model**
- ▶ The general problem involves taking into account **time** and **consequences**
- ▶ Decisions affect the **reward**, the **agent state**, and **environment state**

# Sample-based reinforcement learning

- ▶ Last lecture:
  - ▶ **Planning** by **dynamic programming** to solve a known MDP
- ▶ This lecture:
  - ▶ **Model-free prediction** to **estimate** values in an **unknown** MDP
  - ▶ **Model-free control** to **optimise** values in an **unknown** MDP
- ▶ Not yet:
  - ▶ Learning policies directly in sequential problems (policy gradients)
  - ▶ Continuous MDPs
  - ▶ Deep reinforcement learning

# Sample-based reinforcement learning

- ▶ We can use experience **samples** to learn without a model
- ▶ We call direct sampling of episodes **Monte Carlo**
- ▶ MC is **model-free**: no knowledge of MDP required, only samples

# Sample-based reinforcement learning

- ▶ Simple example, **multi-armed bandit**:
  - ▶ For each action, average reward samples

$$q_t(a) = \frac{\sum_{i=0}^t \mathcal{I}(A_i = a) R_{i+1}}{\sum_{i=0}^t \mathcal{I}(A_i = a)} \approx \mathbb{E}[R_{t+1} | A_t = a] = q(a)$$

- ▶ Equivalently:

$$q_t(A_t) = q_{t-1}(A_t) + \alpha_t (R_t - q_{t-1}(A_t))$$

$$q_t(a) = q_{t-1}(a) \quad \forall a \neq A_t$$

$$\text{with } \alpha_t = \frac{1}{N_t(A_t)} = \frac{1}{\sum_{i=0}^t \mathcal{I}(A_i = a)}$$

Note: we changed notation from  $A_t \rightarrow R_t$  to  $A_t \rightarrow R_{t+1}$

In MDPs, the reward is said to arrive on the time step after the action

## Contextual bandits

- ▶ Consider bandits with different states ('context')
  - ▶ episodes still end after one step
  - ▶ actions do not affect the states
  - ▶ e.g., different visitors to a website
- ▶ Then, we want to estimate

$$q(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

- ▶  $q$  could be a parametric function, e.g., neural network, and we could use loss

$$l_t(\theta) = \frac{1}{2}(q_{\theta}(S_t, A_t) - R_{t+1})^2$$

- ▶ Also works for large (continuous) state spaces  $\mathcal{S}$  — this is just **regression**

# Monte-Carlo Policy Evaluation

- ▶ Now consider sequential decision problems
- ▶ Goal: learn  $v_\pi$  from episodes of experience under policy  $\pi$

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- ▶ The **return** is the total discounted reward (for an episode ending at time  $T > t$ ):

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$$

- ▶ The value function is the expected return:

$$v_\pi(s) = \mathbb{E}[G_t \mid S_t = s, \pi]$$

- ▶ We can just use **sample average** return instead of **expected** return
- ▶ We call this **Monte Carlo policy evaluation**



## Blackjack Example


- ▶ States (200 of them):
  - ▶ Current sum (12-21)
  - ▶ Dealer's showing card (ace-10)
  - ▶ Do I have a "useable" ace? (yes-no)
- ▶ Action **stick**: Stop receiving cards (and terminate)
- ▶ Action **draw**: Take another card (random, no replacement)
- ▶ Reward for **stick**:
  - ▶ +1 if sum of cards  $>$  sum of dealer cards
  - ▶ 0 if sum of cards = sum of dealer cards
  - ▶ -1 if sum of cards  $<$  sum of dealer cards
- ▶ Reward for **draw**:
  - ▶ -1 if sum of cards  $>$  21 (and terminate)
  - ▶ 0 otherwise
- ▶ Transitions: automatically **draw** if sum of cards  $<$  12

# Blackjack Value Function after Monte-Carlo Learning


After 10,000 episodes

After 500,000 episodes



Usable  
ace



+1  
-1



No  
usable  
ace



# Temporal Difference Learning by Sampling Bellman Equations

- ▶ Previous lecture: Bellman equations,

$$v_{\pi}(s) = \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s, A_t \sim \pi(S_t)]$$

- ▶ Previous lecture: Approximate by iterating,

$$v_{k+1}(s) = \mathbb{E} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t \sim \pi(S_t)]$$

- ▶ We can sample this!

$$v_{t+1}(S_t) = R_{t+1} + \gamma v_t(S_{t+1})$$

- ▶ This is likely quite noisy — better to average:

$$v_{t+1}(S_t) = v_t(S_t) + \alpha_t \left( \underbrace{R_{t+1} + \gamma v_t(S_{t+1})}_{\text{target}} - v_t(S_t) \right)$$

# Temporal difference learning

- ▶ In (approximate) DP: we use one step of the model and **bootstrap**

$$\text{target} = \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, \pi]$$

- ▶ In Monte Carlo, we **sample**, and use:

$$\text{target} = G_t (= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + R_T)$$

- ▶ Alternatively, we could **sample and bootstrap**, and use

$$\text{target} = R_{t+1} + \gamma v_t(S_{t+1}).$$

- ▶ This is called **temporal-difference learning**

# Temporal-Difference Learning

- ▶ TD is **model-free** (no knowledge of MDP) and learn directly from experience
- ▶ TD can learn from **incomplete** episodes, by **bootstrapping**
- ▶ TD can learn **during** each episode

# MC and TD

- ▶ Goal: learn  $v_\pi$  online from experience under policy  $\pi$
- ▶ Incremental Monte-Carlo
  - ▶ Update value  $v(S_t)$  towards sampled return  $G_t$

$$v(S_t) \leftarrow v(S_t) + \alpha (G_t - v(S_t))$$

- ▶ Simplest temporal-difference learning algorithm: TD(0)
  - ▶ Update value  $v(S_t)$  towards **estimated** return  $R_{t+1} + \gamma v(S_{t+1})$

$$v(S_t) \leftarrow v(S_t) + \alpha \left( \underbrace{R_{t+1} + \gamma v(S_{t+1})}_{\text{target}} - v(S_t) \right)$$


- ▶  $\delta_t = R_{t+1} + \gamma v(S_{t+1}) - v(S_t)$  is called the **TD error**

## Driving Home Example


<b>State</b>	<b>Elapsed Time (minutes)</b>	<b>Predicted Time to Go</b>	<b>Predicted Total Time</b>
leaving office	0	30	30
reach car, raining	5	35	40
exit highway	20	15	35
behind truck	30	10	40
home street	40	3	43
arrive home	43	0	43

## Driving Home Example: MC vs. TD

Changes recommended by  
Monte Carlo methods ( $\alpha=1$ )



Changes recommended  
by TD methods ( $\alpha=1$ )







# Advantages and Disadvantages of MC vs. TD

- ▶ TD can learn **before** knowing the final outcome
  - ▶ TD can learn online after every step
  - ▶ MC must wait until end of episode before return is known
- ▶ TD can learn **without** the final outcome
  - ▶ TD can learn from incomplete sequences
  - ▶ MC can only learn from complete sequences
  - ▶ TD works in continuing (non-terminating) environments
  - ▶ MC only works for episodic (terminating) environments

## Bias/Variance Trade-Off


- ▶ Return  $G_t = R_{t+1} + \gamma R_{t+2} + \dots$  is an **unbiased** estimate of  $v_\pi(S_t)$
- ▶ TD target  $R_{t+1} + \gamma v(S_{t+1})$  is a **biased** estimate of  $v_\pi(S_t)$ 
  - ▶ Unless  $v(S_{t+1}) = v_\pi(S_{t+1})$
- ▶ But the TD target has much lower variance:
  - ▶ Return depends on **many** random actions, transitions, rewards
  - ▶ TD target depends on **one** random action, transition, reward

## Random Walk Example

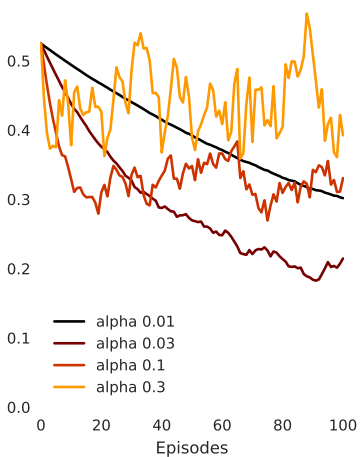


# Random Walk: MC vs. TD

TD



MC



## Batch MC and TD

- ▶ Tabular MC and TD converge:  $v \rightarrow v_\pi$  as experience  $\rightarrow \infty$  and  $\alpha \rightarrow 0$
- ▶ But what about finite experience?

$$\begin{array}{ll} \text{episode 1:} & S_1^1, A_1^1, R_2^1, \dots, S_{T_1}^1 \\ & \vdots \\ \text{episode K:} & S_1^K, A_1^K, R_2^K, \dots, S_{T_K}^K \end{array}$$

- ▶ Repeatedly sample each episodes  $k \in [1, K]$  and apply MC or TD(0)
- ▶ = sampling from an **empirical model**

## AB Example

Two states  $A, B$ ; no discounting; 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$

What is  $v(A), v(B)$ ?

## AB Example

Two states  $A, B$ ; no discounting; 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$


$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$



What is  $v(A)$ ,  $v(B)$ ?

## Certainty Equivalence

- ▶ MC converges to best mean-squared fit for the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} \left( G_t^k - v(S_t^k) \right)^2$$

- ▶ In the AB example,  $v(A) = 0$
- ▶ TD converges to solution of max likelihood Markov model
  - ▶ Solution to the empirical MDP  $\langle \mathcal{S}, \mathcal{A}, \hat{p}, \gamma \rangle$  that best fits the data
  - ▶ In the AB example:  $\hat{p}(S_{t+1} = B \mid S_t = A) = 1$ , and therefore  $v(A) = v(B) = 0.75$




## Advantages and Disadvantages of MC vs. TD (3)

- ▶ TD exploits Markov property
  - ▶ Usually more efficient in Markov environments
- ▶ MC does not exploit Markov property
  - ▶ Usually more accurate in non-Markov environments


# Monte-Carlo Backup

$$v(S_t) \leftarrow v(S_t) + \alpha (G_t - v(S_t))$$




# Temporal-Difference Backup

$$v(S_t) \leftarrow v(S_t) + \alpha (R_{t+1} + \gamma v(S_{t+1}) - v(S_t))$$



# Dynamic Programming Backup


$$v(S_t) \leftarrow \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid A_t \sim \pi(S_t)]$$



# Bootstrapping and Sampling

- ▶ **Bootstrapping**: update involves an estimate
  - ▶ MC does not bootstrap
  - ▶ DP bootstraps
  - ▶ TD bootstraps
- ▶ **Sampling**: update samples an expectation
  - ▶ MC samples
  - ▶ DP does not sample
  - ▶ TD samples


# Unified View of Reinforcement Learning



# Model-Free Control


- ▶ Previous: **Model-free prediction:**  
**Estimate** the value function of an unknown MDP
- ▶ Next: **Model-free control:**  
**Optimise** the value function of an unknown MDP

# Generalized Policy Iteration (Refresher)



**Policy evaluation** Estimate  $v_\pi$   
e.g. Iterative policy evaluation

**Policy improvement** Generate  $\pi' \geq \pi$   
e.g. Greedy policy improvement





# Monte Carlo

- ▶ Recall, Monte Carlo estimate from state  $S_t$  is

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$

- ▶  $\mathbb{E}[G_t] = v_\pi$
- ▶ So, we can average multiple estimates to get  $v_\pi$

# Model-Free Policy Iteration Using Action-Value Function


- ▶ But, greedy policy improvement over  $v(s)$  requires model of MDP

$$\pi'(s) = \operatorname{argmax}_a \mathbb{E} [R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s, A_t = a]$$

- ▶ Greedy policy improvement over  $q(s, a)$  is model-free

$$\pi'(s) = \operatorname{argmax}_a q(s, a)$$


# Generalised Policy Iteration with Action-Value Function



**Policy evaluation** Monte-Carlo policy evaluation,  $q \approx q_\pi$

**Policy improvement** Greedy policy improvement? No exploration!

## Monte-Carlo Generalized Policy Iteration



Every episode:

Policy evaluation Monte-Carlo policy evaluation,  $q \approx q_\pi$

Policy improvement  $\epsilon$ -greedy policy improvement

# GLIE

## Definition

### Greedy in the Limit with Infinite Exploration (GLIE)

- ▶ All state-action pairs are explored infinitely many times,

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- ▶ The policy converges to a greedy policy,

$$\lim_{k \rightarrow \infty} \pi_k(s, a) = \mathcal{I}(a = \operatorname{argmax}_{a'} q_k(s, a'))$$

- ▶ For example,  $\epsilon$ -greedy with  $\epsilon_k = \frac{1}{k}$

## GLIE Monte-Carlo Control

- ▶ Sample  $k$ th episode using  $\pi$ :  $\{S_1, A_1, R_2, \dots, S_T\} \sim \pi$
- ▶ For each state  $S_t$  and action  $A_t$  in the episode,

$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$

$$q(S_t, A_t) \leftarrow q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - q(S_t, A_t))$$

- ▶ Improve policy based on new action-value function

$$\epsilon \leftarrow 1/k$$

$$\pi \leftarrow \epsilon\text{-greedy}(q)$$


### Theorem

*GLIE Monte-Carlo control converges to the optimal action-value function,  
 $q(s, a) \rightarrow q_*(s, a)$*

# MC vs. TD Control

- ▶ Temporal-difference (TD) learning has several advantages over Monte-Carlo (MC)
  - ▶ Lower variance
  - ▶ Online
  - ▶ Can learn from incomplete sequences
- ▶ Natural idea: use TD instead of MC for control
  - ▶ Apply TD to  $q(s, a)$
  - ▶ Use, e.g.,  $\epsilon$ -greedy policy improvement
  - ▶ Update every time-step


## Updating Action-Value Functions with Sarsa



$$q(s, a) \leftarrow q(s, a) + \alpha (r + \gamma q(s', a') - q(s, a))$$



# Sarsa



Every **time-step**:

Policy evaluation **Sarsa**,  $q \approx q_\pi$

Policy improvement  $\epsilon$ -greedy policy improvement

## Tabular Sarsa

Initialize  $Q(s, a)$  arbitrarily

Repeat (for each episode):

    Initialize  $s$

    Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

    Repeat (for each step of episode):

        Take action  $a$ , observe  $r, s'$

        Choose  $a'$  from  $s'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$

$s \leftarrow s'; a \leftarrow a';$

    until  $s$  is terminal

# On and Off-Policy Learning

- ▶ **On-policy** learning
  - ▶ “Learn on the job”
  - ▶ Learn about policy  $\pi$  from experience sampled from  $\pi$
- ▶ **Off-policy** learning
  - ▶ “Look over someone’s shoulder”
  - ▶ Learn about policy  $\pi$  from experience sampled from  $b$

# Dynamic programming

- We discussed other algorithms:

$$v_{k+1}(s) = \mathbb{E} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t \sim \pi(S_t)] \quad (\text{policy evaluation})$$

$$v_{k+1}(s) = \max_a \mathbb{E} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t = a] \quad (\text{value iteration})$$

$$q_{k+1}(s, a) = \mathbb{E} [R_{t+1} + \gamma q_k(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a] \quad (\text{policy evaluation})$$

$$q_{k+1}(s, a) = \mathbb{E} \left[ R_{t+1} + \gamma \max_{a'} q_k(S_{t+1}, a') \mid S_t = s, A_t = a \right] \quad (\text{value iteration})$$

# TD learning

- ▶ Analogous TD algorithms

$$v_{t+1}(S_t) = v_t(S_t) + \alpha_t (R_{t+1} + \gamma v_t(S_{t+1}) - v_t(S_t)) \quad (\text{TD})$$

$$q_{t+1}(s, a) = q_t(S_t, A_t) + \alpha_t (R_{t+1} + \gamma q_t(S_{t+1}, A_{t+1}) - q_t(S_t, A_t)) \quad (\text{Sarsa})$$

$$q_{t+1}(s, a) = q_t(S_t, A_t) + \alpha_t \left( R_{t+1} + \gamma \max_{a'} q_t(S_{t+1}, a') - q_t(S_t, A_t) \right) \quad (\text{Q-learning})$$

- ▶ Note, no trivial analogous version of value iteration

$$v_{k+1}(s) = \max_a \mathbb{E} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t = a]$$

Can you explain why? We do not have a model and can not sample.

# Off-Policy Learning

- ▶ Evaluate target policy  $\pi(s, a)$  to compute  $v_\pi(s)$  or  $q_\pi(s, a)$
- ▶ While following behaviour policy  $b(s, a)$

$$\{S_1, A_1, R_2, \dots, S_T\} \sim b$$

- ▶ Why is this important?
  - ▶ Learn from observing humans or other agents
  - ▶ Re-use experience from old policies
  - ▶ Learn about **optimal** policy while following **exploratory** policy
  - ▶ Learn about **multiple** policies while following **one** policy
- ▶ **Q-learning** estimates the value of the **greedy** policy

$$q_{t+1}(s, a) = q_t(S_t, A_t) + \alpha_t \left( R_{t+1} + \gamma \max_{a'} q_t(S_{t+1}, a') - q_t(S_t, A_t) \right)$$



# Q-Learning Control Algorithm

## Theorem

*Q-learning control converges to the optimal action-value function,  $q \rightarrow q^*$ , as long as we take each action in each state infinitely often.*

Note: no need for greedy behaviour!

# Cliff Walking Example





# Q-learning overestimation

- ▶ Classical Q-learning has potential issues
- ▶ Recall

$$\max_a q_t(S_{t+1}, a) = q_t(S_{t+1}, \operatorname{argmax}_a q_t(S_{t+1}, a))$$

- ▶ Uses same values to **select** and to **evaluate**
- ▶ ... but values are approximate
  - ▶ more likely to select **overestimated values**
  - ▶ less likely to select **underestimated values**
- ▶ This causes upward bias

# Double Q-learning

- ▶ Q-learning uses same values to **select** and to **evaluate**

$$R_{t+1} + \gamma q_t(S_{t+1}, \underset{a}{\operatorname{argmax}} q_t(S_{t+1}, a))$$



- ▶ Solution: decouple selection from evaluation, using the
- ▶ **Double Q-learning:**
  - ▶ Store two  $q$  functions:  $q, q'$

$$R_{t+1} + \gamma \underset{a}{\operatorname{argmax}} q'_t(S_{t+1}, a) q_t(S_{t+1}, a) \quad (1)$$

$$R_{t+1} + \gamma q_t(S_{t+1}, \underset{a}{\operatorname{argmax}} q'_t(S_{t+1}, a)) \quad (2)$$

- ▶ Each step, pick one (e.g., randomly) and update, using update (1) for  $q$  or (2) for  $q'$
- ▶ Can use both to act (e.g., use  $q + q'$ )


## Roulette example



# Double DQN on Atari

DQN

Double DQN



# Importance Sampling

- ▶ Estimate the expectation of a different distribution

$$\begin{aligned}\mathbb{E}_{x \sim d}[f(x)] &= \sum d(x)f(x) \\ &= \sum d'(x) \frac{d(x)}{d'(x)} f(x) \\ &= \mathbb{E}_{x \sim d'} \left[ \frac{d(x)}{d'(x)} f(x) \right]\end{aligned}$$

# Importance Sampling

- ▶ Estimate the expectation of a different distribution

$$\begin{aligned}\mathbb{E}[R_{t+1} \mid S_t = s, A_t \sim \pi] &= \sum_a \pi(a|s) r(s, a) \\ &= \sum_a b(a|s) \frac{\pi(a|s)}{b(a|s)} r(s, a) \\ &= \mathbb{E} \left[ \frac{\pi(A_t|S_t)}{b(A_t|S_t)} R_{t+1} \mid S_t = s, A_t \sim b \right]\end{aligned}$$

- ▶ Ergo, when following  $b$ , can use  $\frac{\pi(A_t|S_t)}{b(A_t|S_t)} R_{t+1}$  as unbiased sample

# Importance Sampling for Off-Policy Monte-Carlo

- ▶ Use returns generated from  $b$  to evaluate  $\pi$
- ▶ Weight return  $G_t$  according to similarity between policies

$$G_t^{\pi/b} = \frac{\pi(S_t, A_t)}{b(S_t, A_t)} \frac{\pi(S_{t+1}, A_{t+1})}{b(S_{t+1}, A_{t+1})} \cdots \frac{\pi(S_{T-1}, A_{T-1})}{b(S_{T-1}, A_{T-1})} G_t$$

- ▶ Update towards **corrected** return

$$v(S_t) \leftarrow v(S_t) + \alpha \left( G_t^{\pi/b} - v(S_t) \right)$$

- ▶  $\mathbb{E} \left[ G_t^{\pi/b} \mid S_t = s, b \right] = v_\pi(s)$  — **no bias!**
- ▶ ...but importance sampling can dramatically **increase variance**...

## Importance Sampling for Off-Policy TD Updates

- ▶ Use TD targets generated from  $b$  to evaluate  $\pi$
- ▶ Weight TD target  $r + \gamma v(s')$  by importance sampling
- ▶ Only need a single importance sampling correction

$$v(S_t) \leftarrow v(S_t) + \alpha \left( \frac{\pi(S_t, A_t)}{b(S_t, A_t)} (R_{t+1} + \gamma v(S_{t+1})) - v(S_t) \right)$$

- ▶ Much lower variance than Monte-Carlo importance sampling
- ▶ Policies only need to be similar over a single step



## Q-Learning

- ▶ We now consider off-policy learning of action-values  $q(s, a)$
- ▶ **No** importance sampling is required
- ▶ Next action may be chosen using behaviour policy  $A_{t+1} \sim b(S_{t+1}, \cdot)$
- ▶ But we consider probabilities under  $\pi(S_t, \cdot)$
- ▶ Update  $q(S_t, A_t)$  towards value of alternative action

$$q(S_t, A_t) \leftarrow q(S_t, A_t) + \alpha \left( R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) q(S_{t+1}, a) - q(S_t, A_t) \right)$$

- ▶ Called **Expected Sarsa** (when  $b = \pi$ ) or **General Q-learning**

# Off-Policy Control with Q-Learning

- ▶ We want behaviour and target policies to **improve**
- ▶ E.g., the target policy  $\pi$  is **greedy** w.r.t.  $q(s, a)$

$$\pi(S_{t+1}) = \operatorname{argmax}_{a'} q(S_{t+1}, a')$$

- ▶ The behaviour policy  $b$  is e.g.  **$\epsilon$ -greedy** w.r.t.  $q(s, a)$
- ▶ The Q-learning target then simplifies:

$$\begin{aligned} R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) q(S_{t+1}, a) \\ = R_{t+1} + \gamma \max_a q(S_{t+1}, a) \end{aligned}$$

## Questions?

*The only stupid question is the one you were afraid to ask but never did.*  
-Rich Sutton

For questions that arise outside of class, please use Moodle